

CS267 Assignment 2

Hussain Al Salem Jason Poulos Yang You

March 1, 2016

1 Introduction

In this report, we describe several parallel implementations of a 2D particle simulator and report their performance. The goal is to parallelize code that runs in time $T = O(n)$ on a single processor to run in time T/p when using p processors by taking advantage of shared and distributed memory models. Specifically, we try three implementations: serial code that runs in $O(n)$ time; a MPI distributed memory implementation that runs in $O(n)$ time and $O(n/p)$ scaling; and a OpenMP shared memory implementation. For Part 2 of the assignment, we implement in GPU.

2 Serial Implementation

2.1 Data structures

Our serial implementation uses a binning structure to execute the interaction step in $O(n)$ time. The 2D grid is divided into small squared bins where each side has the size of cutoff. The size of cutoff was used since it is the maximum distance of interaction. Thus, each particle can only interact with particles in its own square and its 8 neighboring squares on a single iteration. Furthermore, bins that are attached to the edges will have less neighbors and thus they were fixed accordingly. The small bins are stored in a data structure type where each bin has a unique consecutive index. Each bin points to the particles indices contained inside the bin. After each iteration, we empty the bins and update them according to the new particle locations. In the interaction step, each particle only need to check for particles within its bin and neighboring bins. Since the number of particles per square is bound (and equal to 1 on average), the interaction step takes $O(n)$ time.

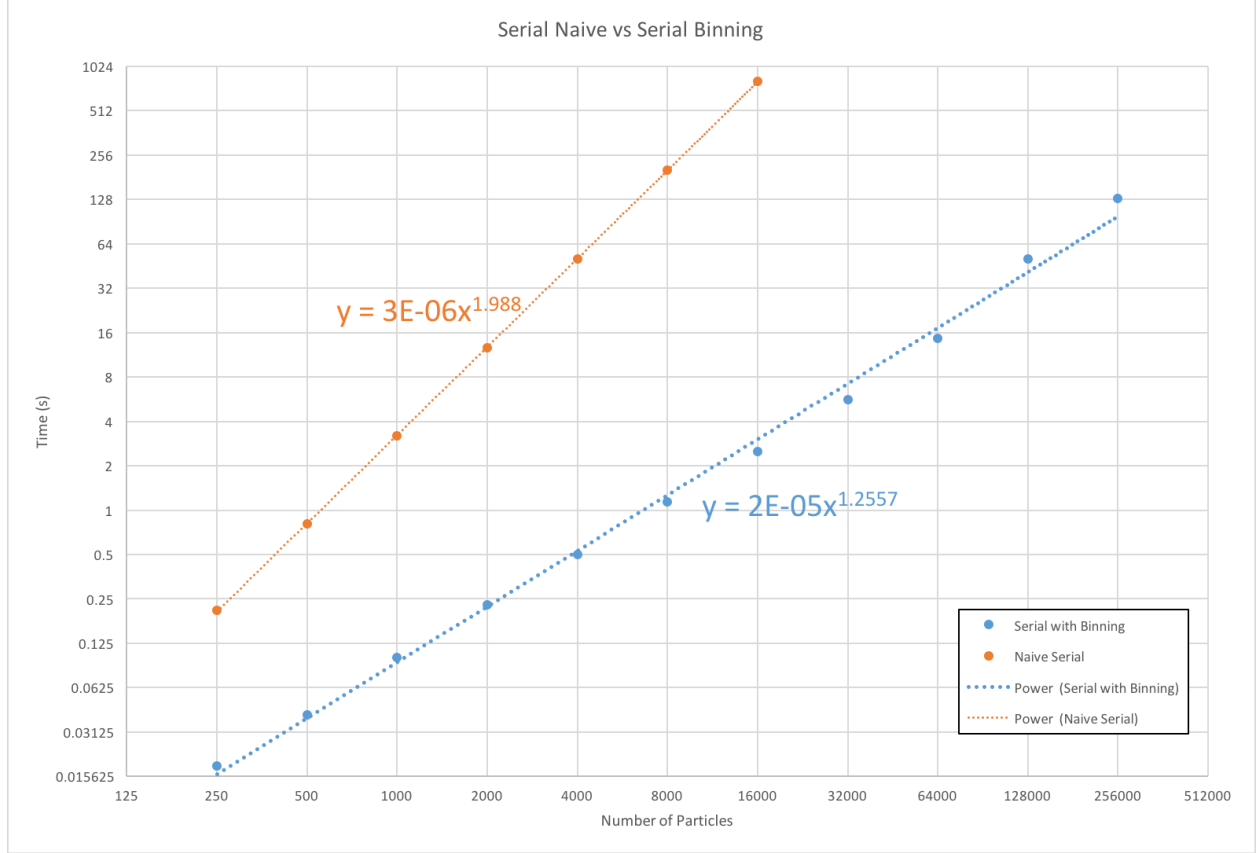


Figure 1: A plot in log-log scale that shows serial without binning runs on $O(n^2)$ time and serial with binning runs in $O(n)$ time.

Figure 2: A plot in log-linear scale that shows performance as a percent of peak performance for different numbers of processors.

2.2 Results

Figure 1 shows that our serial binning algorithm follows linear scaling ($O(n)$). For comparison, it shows the performance of the naive $O(n^2)$ algorithm on smaller problems.

3 OpenMP implementation

3.1 Synchronization

A description of the synchronization you used in the shared memory implementation.

Figure 3: $O(n)$ scaling of the OpenMP implementation.

Figure 4: Weak scaling of the OpenMP implementation.

3.2 Results

4 MPI Implementation

4.1 Communication between nodes

A description of the communication you used in the distributed memory implementation.

4.2 Results

5 Comparison of distributed and shared implementations

5.1 Where does the time go?

Consider breaking down the runtime into computation time, synchronization time and/or communication time. How do they scale with p ?

5.2 Discussion

A description of the design choices that you tried and how did they affect the performance.

A discussion on whether it is possible to do better

A discussion on using pthreads, OpenMP and MPI.

6 GPU Implementation

A description of any synchronization needed A description of any GPU-specific optimizations you tried

Figure 5: Strong scaling of the OpenMP implementation.

Figure 6: $O(n)$ scaling of the MPI implementation.

Figure 7: Weak scaling of the MPI implementation.

6.1 Results

6.2 Discussion

A discussion on the strengths and weaknesses of CUDA and the current GPU architecture

Figure 8: Strong scaling of the MPI implementation.

Figure 9: Speedup plots that show how closely parallel codes approach the idealized p -times speedup.

Figure 10: A plot of the speedup of the GPU code versus the serial, openmp, mpi runs on the CPU of the node

Figure 11: A plot in log-log scale that shows the performance of your code versus the naive GPU code.