

# CS267 Assignment 2

Hussain Al Salem    Jason Poulos    Yang You

February 20, 2016

## 1 Introduction

In this report, we describe several parallel implementations of a 2D particle simulator and report their performance. The goal is to parallelize code that runs in time  $T = O(n)$  on a single processor to run in time  $T/p$  when using  $p$  processors by taking advantage of shared and distributed memory models. Specifically, we try three implementations: serial code that runs in  $O(n)$  time; a MPI distributed memory implementation that runs in  $O(n)$  time and  $O(n/p)$  scaling; and a OpenMP shared memory implementation. For Part 2 of the assignment, we implement in GPU.

## 2 Serial Implementation

### 2.1 Data structures

Describe the serial structures used to achieve  $O(n)$ .

### 2.2 Results

## 3 OpenMP implementation

### 3.1 Synchronization

A description of the synchronization you used in the shared memory implementation.

Figure 1: A plot in log-log scale that shows that serial and parallel codes run in  $O(n)$  time.

Figure 2: A plot in log-linear scale that shows performance as a percent of peak performance for different numbers of processors.

Figure 3:  $O(n)$  scaling of the OpenMP implementation.

## 3.2 Results

# 4 MPI Implementation

## 4.1 Communication between nodes

A description of the communication you used in the distributed memory implementation.

## 4.2 Results

# 5 Comparison of distributed and shared implementations

## 5.1 Where does the time go?

Consider breaking down the runtime into computation time, synchronization time and/or communication time. How do they scale with  $p$ ?

## 5.2 Discussion

A description of the design choices that you tried and how did they affect the performance.

A discussion on whether it is possible to do better

A discussion on using pthreads, OpenMP and MPI.

# 6 GPU Implementation

A description of any synchronization needed A description of any GPU-specific optimizations you tried

Figure 4: Weak scaling of the OpenMP implementation.

Figure 5: Strong scaling of the OpenMP implementation.

Figure 6:  $O(n)$  scaling of the MPI implementation.

## 6.1 Results

## 6.2 Discussion

A discussion on the strengths and weaknesses of CUDA and the current GPU architecture

Figure 7: Weak scaling of the MPI implementation.

Figure 8: Strong scaling of the MPI implementation.

Figure 9: Speedup plots that show how closely parallel codes approach the idealized  $p$ -times speedup.

Figure 10: A plot of the speedup of the GPU code versus the serial, openmp, mpi runs on the CPU of the node

Figure 11: A plot in log-log scale that shows the performance of your code versus the naive GPU code.